Here is your document with **100% original content preserved** — no rewording, no omissions, no restructuring of meaning.
Only **spacing, line breaks, and visual hierarchy** were added for readability.

Claude Code for Product Teams — Workshop Syllabus

> **A 4-hour, hands-on Claude Code workshop for Product Managers and Product Designers.**
>
> **Next cohort:** Thursday, May 7, 2026 · 17:00 – 21:00 · [LOCATION]
>
> Designed and delivered by Shmulik Davar / BrAIght Wave.
>
> Format: in-person, 40 participants, 10 pods of 4. Mixed levels per pod.

Why this workshop exists

Most AI tooling for product teams is still treated as a chat window — copy a question in, copy an answer out.

Claude Code changes that mental model: it's an AI agent that runs on the participant's own laptop, reads their real files, edits real code, and connects to real tools (Jira, Notion, Figma, GitHub) through MCP.

This workshop teaches Product Managers and Product Designers how to use Claude Code as a working tool — not to become engineers, but to do their existing jobs (scoping, PRDs, design-system enforcement, customer feedback synthesis) at higher leverage.

By 16:00, every participant walks out with three working **skills** saved in their `.claude/skills/` folder, framed as the first three entries of a team Claude Code library.

They've also seen four practitioners share what real production use looks like, including one explicitly contrarian voice.

Audience and prerequisites

Audience: Working PMs and product designers, mixed levels:

- Some have never used Claude Code (or any agentic AI tool).
- Some have used Cursor, Lovable, Bolt, or v0 for prototyping.
- A small subset are advanced Claude Code users.

We assemble pods of 4 with **1 advanced + 2 mid + 1 beginner per pod** based on a pre-workshop skill survey. Peer teaching is a load-bearing element of the design.

Prerequisites (sent 1 week before):

- A laptop (macOS 13+, Windows 10+, or Linux Ubuntu 20.04+).
- A paid Claude account (Pro $20/mo recommended; Max, Teams, or Enterprise also work). Free tier does NOT include Claude Code.
- Claude Desktop app installed and signed in.
- Git installed and configured (name + email).
- A GitHub account.
- A Figma account with edit access to at least one file.

A separate setup guide walks participants through all of this and includes a verification step (post a screenshot in the WhatsApp room) that runs 5 days before the workshop.

By the time anyone walks in, they're green or flagged.

Surface choice — why we use the Claude Desktop app

Claude Code ships in three flavors: Desktop app, CLI (terminal), and IDE extensions.

We use the **Desktop app** for the entire workshop because:

- The audience is mostly beginners. Adding "learn the terminal" on top of "learn agentic AI" doubles the cognitive load.
- Every pattern we teach (plan mode, CLAUDE.md, skills, MCP) works identically across all three flavors.
- Skills saved during the workshop are portable — participants who graduate to the CLI later don't lose anything.
- Demos are easier to follow visually: the file tree, plan sidebar, and split-pane views show what Claude is doing in real time.

The CLI and IDE extensions are mentioned in the keynote and cheat sheet as “what advanced participants might use after the workshop.” The patterns transfer 1:1.

Workshop goals

By the end of the session, every participant can:

1. Use Claude Code to **explore an unfamiliar codebase and scope a feature** with file-level evidence — not vibes.
2. Use Claude Code with **Figma + a design system** to translate design into working code that respects system tokens, not freelance pixels.
3. **Synthesize real customer feedback** into a feature definition, then run a **PRD round-trip** — PM ↔ Dev ↔ PM — entirely inside Claude Code.
4. Place Claude Code correctly in their stack (vs. Claude.ai, Cursor, Lovable, v0, Codex).
5. Walk out with **3 skills, 1 cheat sheet, and 1 commitment for Friday**.

Format and structure

17:00–17:10 Pod intros (10 min)
17:10–17:40 Keynote: Why Claude Code, why now (30 min)
17:40–18:30 PART 1: EXPLORE — Onboard to a new product → save as skill (50 min)
18:30–18:40 ☕ Break
18:40–19:30 PART 2: CONSTRAIN — Figma + Design System → save as skill (50 min)
19:30–19:40 ☕ Break
19:40–20:30 PART 3: COLLABORATE — Customer feedback → PRD round-trip → save as skill (50 min)
20:30–20:50 Graduate panel: 4 voices from the field (20 min)
20:50–21:00 Closing: What's coming this year (10 min)

Pedagogical principle: every Part follows the same shape — short demo (10–15 min) → independent pod work (~25 min) with the instructor walking the room → share-back from 2-3 pods → save the working pattern as a `SKILL.md` file.

Three Parts means three skills saved by 21:00.

Pod rules (introduced in the keynote):

1. Rotate the keyboard. Different person types each Part.
2. Beginner reads Claude’s output aloud — this is how the room catches Claude when it’s confidently wrong.
3. Push back on Claude. The first answer is rarely the best one.
4. Ask your pod before asking the instructor. Pedagogy happens at the table.
5. Floaters handle setup problems only. Pedagogy is the pod’s job.